

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
```

```
df = pd.read_csv('cyberbullying.csv.zip')
```

```
df.head()
```



	tweet_text	cyberbullying_type
0	In other words #katandandre, your food was cra...	not_cyberbullying
1	Why is #aussietv so white? #MKR #theblock #ImA...	not_cyberbullying
2	@XochitlSuckkks a classy whore? Or more red ve...	not_cyberbullying
3	@Jason_Gio meh. :P thanks for the heads up, b...	not_cyberbullying
4	@RudhoeEnglish This is an ISIS account pretend...	not_cyberbullying

[+ Code](#)
[+ Text](#)

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 47692 entries, 0 to 47691
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   tweet_text      47692 non-null  object
1   cyberbullying_type 47692 non-null  object
dtypes: object(2)
memory usage: 745.3+ KB
```

```
!pip install demoji
```

```
Collecting demoji
  Downloading demoji-1.1.0-py3-none-any.whl (42 kB)
    42.9/42.9 kB 1.2 MB/s eta 0:00:00
Installing collected packages: demoji
Successfully installed demoji-1.1.0
```

```
import re
from nltk.corpus import stopwords
from nltk.stem.snowball import SnowballStemmer
import demoji
import string
```

```
import warnings
warnings.filterwarnings("ignore")
from warnings import simplefilter
from sklearn.exceptions import ConvergenceWarning
simplefilter("ignore", category=ConvergenceWarning)
```

```
import nltk
nltk.download('stopwords')
STOPWORDS = set(stopwords.words('english'))
STOPWORDS.update(['rt', 'mkr', 'didn', 'bc', 'n', 'm',
                  'im', 'll', 'y', 've', 'u', 'un', 'don',
                  'p', 't', 's', 'aren', 'kp', 'o', 'kat',
                  'de', 're', 'amp', 'will', 'wa', 'e', 'like'])
stemmer = SnowballStemmer('english')
def clean_text(text):
    pattern = re.compile(r"#[A-Za-z0-9]+|@[A-Za-z0-9]+|https?://\S+|www\.\S+|\S+\.[a-z]+|RT @")
    text = pattern.sub('', text)
    text = " ".join(text.split())
    text = text.lower()
    text = " ".join([stemmer.stem(word) for word in text.split()])
    remove_punc = re.compile(r"%s" % re.escape(string.punctuation))
    text = remove_punc.sub('', text)
    text = " ".join([word for word in str(text).split() if word not in STOPWORDS])
    emoji = demoji.findall(text)
    for emot in emoji:
        text = re.sub(r"(%s)" % (emot), "_" .join(emoji[emot].split()), text)

    return text

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.

df['cleaned_text'] = df['tweet_text'].apply(lambda text: clean_text(text))

df.head()

    tweet_text  cyberbullying_type  cleaned_text
0  In other words #katandandre, your food was cra...  not_cyberbullying  word food crapilicious
1  Why is #aussietv so white? #MKR #theblock #ImA...  not_cyberbullying  whi white
2  @XochitlSuckkks a classy whore? Or more red ve...  not_cyberbullying  classi whore red velvet cupcakes
3  @Jason_Gio meh. :P thanks for the heads up, b...  not_cyberbullying  gio meh thank head concern anoth angri dude tw...
4  @RudhoeEnglish This is an ISIS account pretend...  not_cyberbullying  isi account pretend kurdish account islam lies

df.isnull().sum()

tweet_text      0
cyberbullying_type  0
cleaned_text    0
dtype: int64

df['cleaned_text'].duplicated().sum()

2887

df.drop_duplicates("cleaned_text", inplace = True)

df['cleaned_text'].str.isspace().sum()

0

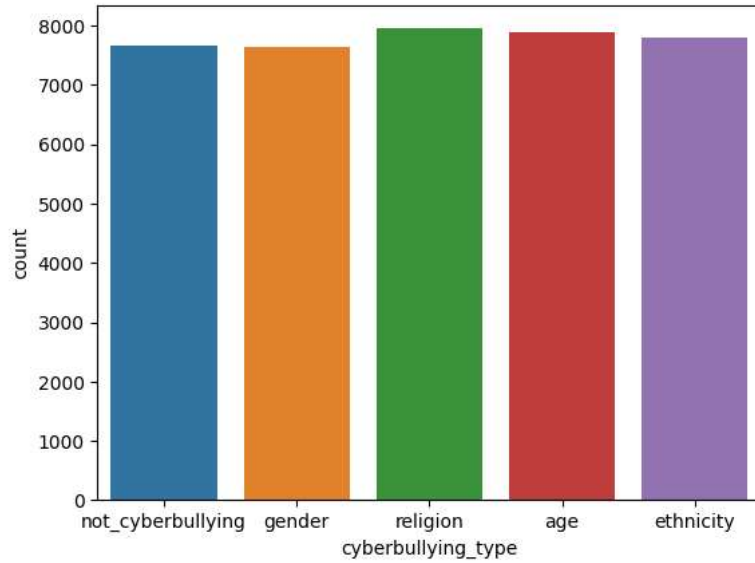
df = df[df["cyberbullying_type"]!="other_cyberbullying"]

df['cyberbullying_type'].value_counts()

religion      7946
age           7887
ethnicity     7797
not_cyberbullying  7670
gender        7637
Name: cyberbullying_type, dtype: int64

sns.countplot(data = df, x = 'cyberbullying_type')
```

<Axes: xlabel='cyberbullying_type', ylabel='count'>



```
for cyber_type in df.cyberbullying_type.unique():
```

```
    top50_word = df.cleaned_text[df.cyberbullying_type==cyber_type].str.split(expand=True).stack().value_counts()[:15]
```

```
    fig = px.bar(top50_word, color=top50_word.values, color_continuous_scale=px.colors.sequential.RdPu, custom_data=[top50_word.values])
```

```
    fig.update_traces(marker_color='red')
```

```
    fig.update_traces(hovertemplate='<b>Count: </b>{customdata[0]}')
```

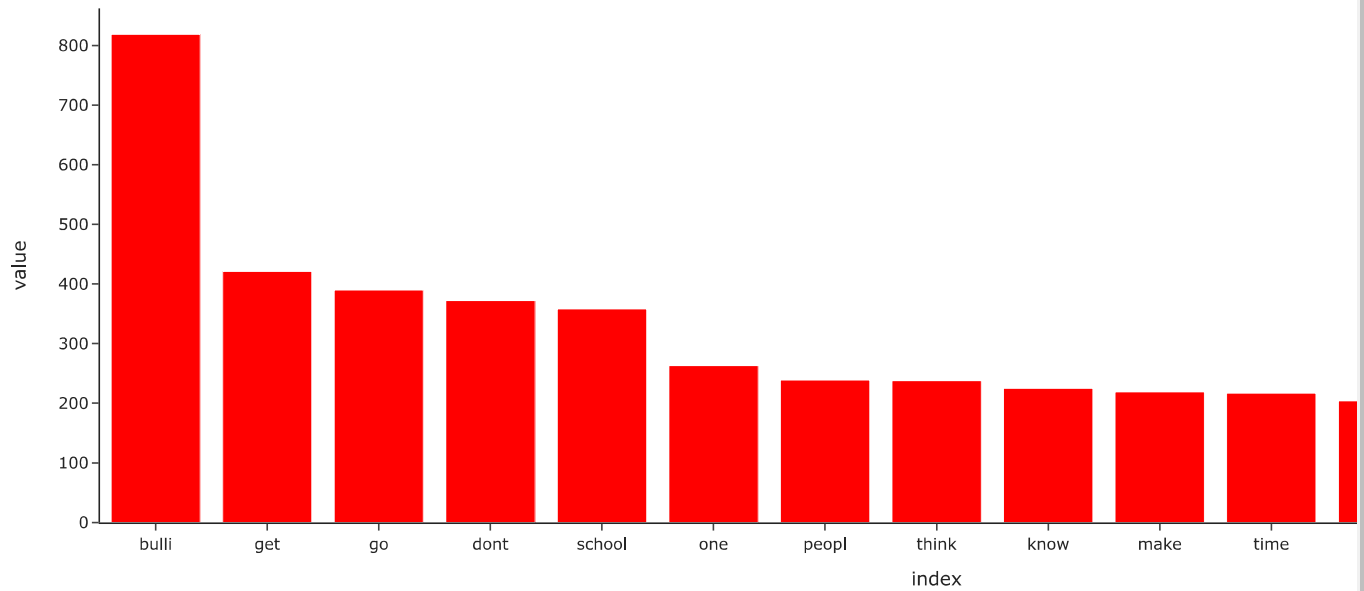
```
    fig.update_layout(title=f"Top 15 words for {cyber_type}",
```

```
                      template='simple_white',
```


```
                      hovermode='x unified')
```

```
    fig.show()
```

Top 15 words for not_cyberbullying



Top 15 words for gender



```
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
X = df['cleaned_text']
y = df['cyberbullying_type']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.1, random_state = 42)
```

```
tfidf = TfidfVectorizer(max_features = 5000)
```

```
X_train_tfidf = tfidf.fit_transform(X_train)
```

```
X_test_tfidf = tfidf.transform(X_test)
```

```
X_train_tfidf
```

```
<35043x5000 sparse matrix of type '<class 'numpy.float64''
with 403374 stored elements in Compressed Sparse Row format>
```

```
X_test_tfidf
```

```
<3894x5000 sparse matrix of type '<class 'numpy.float64'>'
  with 44224 stored elements in Compressed Sparse Row format>
```

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
tfidf_array_train = X_train_tfidf.toarray()
tfidf_array_test = X_test_tfidf.toarray()
scaled_X_train = scaler.fit_transform(tfidf_array_train)
scaled_X_test = scaler.transform(tfidf_array_test)
```

```
8000+
```

```
from sklearn.decomposition import PCA
NUM_COMPONENTS = 5000
pca = PCA(NUM_COMPONENTS)
reduced = pca.fit(scaled_X_train)
```

```
|
```


```
variance_explained = np.cumsum(pca.explained_variance_ratio_)
```

```
=
```

```
|
```

```
fig, ax = plt.subplots(figsize=(8, 6))
plt.plot(range(NUM_COMPONENTS), variance_explained, color='r')
ax.grid(True)
plt.xlabel("Number of components")
plt.ylabel("Cumulative explained variance")
```

```
Text(0, 0.5, 'Cumulative explained variance')
```



```
|
```

```
final_pca = PCA(0.9)
reduced_90 = final_pca.fit_transform(scaled_X_train)
```

```
index
```

```
reduced_90_test = final_pca.transform(scaled_X_test)
```

```
reduced_90.shape
```

```
(35043, 3999)
```

```
final_pca = PCA(0.8)
reduced_80 = final_pca.fit_transform(scaled_X_train)
```

```
reduced_80.shape
```

```
(35043, 3290)
```

```
from sklearn.metrics import confusion_matrix, classification_report
```

```

from sklearn.linear_model import LogisticRegression
log_model_pca = LogisticRegression()
log_model_pca.fit(reduced_90, y_train)
preds_log_model_pca = log_model_pca.predict(reduced_90_test)
print(classification_report(y_test, preds_log_model_pca))
confusion_matrix(y_test, preds_log_model_pca)

```

	precision	recall	f1-score	support
age	0.86	0.83	0.85	766
ethnicity	0.90	0.85	0.87	801
gender	0.77	0.79	0.78	788
not_cyberbullying	0.65	0.67	0.66	783
religion	0.84	0.86	0.85	756
accuracy			0.80	3894
macro avg	0.80	0.80	0.80	3894
weighted avg	0.80	0.80	0.80	3894

```

array([[638, 10, 28, 76, 14],
       [ 19, 682, 24, 50, 26],
       [ 20, 19, 622, 100, 27],
       [ 47, 30, 119, 528, 59],
       [ 14, 18, 14, 60, 650]])

```

```

from sklearn.experimental import enable_halving_search_cv
from sklearn.model_selection import HalvingGridSearchCV
log_model = LogisticRegression(solver = 'saga')
param_grid = {'C': np.logspace(0, 10, 5)}
grid_log_model = HalvingGridSearchCV(log_model, param_grid = param_grid, n_jobs = -1, min_resources = 'exhaust', factor = 3)
grid_log_model.fit(X_train_tfidf, y_train)
preds_grid_log_model = grid_log_model.predict(X_test_tfidf)
print(classification_report(y_test, preds_grid_log_model))
confusion_matrix(y_test, preds_grid_log_model)

```

	precision	recall	f1-score	support
age	0.96	0.97	0.96	766
ethnicity	0.98	0.98	0.98	801
gender	0.92	0.84	0.88	788
not_cyberbullying	0.80	0.85	0.82	783
religion	0.94	0.96	0.95	756
accuracy			0.92	3894
macro avg	0.92	0.92	0.92	3894
weighted avg	0.92	0.92	0.92	3894

```

array([[743, 1, 3, 18, 1],
       [ 2, 782, 2, 13, 2],
       [ 1, 6, 660, 113, 8],
       [ 28, 7, 49, 666, 33],
       [ 1, 2, 2, 25, 726]])

```

```
grid_log_model.best_estimator_
```

```

LogisticRegression
LogisticRegression(solver='saga')

```

```

from sklearn.svm import LinearSVC
svm_model = LinearSVC()
C = [1e-5, 1e-4, 1e-2, 1e-1, 1]
param_grid = {'C': C}
grid_svm_model = HalvingGridSearchCV(svm_model, param_grid = param_grid, n_jobs = -1, min_resources = 'exhaust', factor = 3)
grid_svm_model.fit(X_train_tfidf, y_train)
preds_grid_svm_model = grid_svm_model.predict(X_test_tfidf)
print(classification_report(y_test, preds_grid_svm_model))
confusion_matrix(y_test, preds_grid_svm_model)

```

	precision	recall	f1-score	support
age	0.94	0.98	0.96	766
ethnicity	0.97	0.98	0.98	801
gender	0.94	0.81	0.87	788
not_cyberbullying	0.79	0.85	0.82	783
religion	0.95	0.96	0.96	756
accuracy			0.92	3894

macro avg	0.92	0.92	0.92	3894
weighted avg	0.92	0.92	0.92	3894

```
array([[754, 1, 1, 10, 0],
       [ 2, 783, 2, 13, 1],
       [ 3, 7, 637, 133, 8],
       [39, 11, 36, 665, 32],
       [ 1, 2, 1, 23, 729]])
```

```
grid_svm_model.best_estimator_
```

```
▼ LinearSVC
LinearSVC(C=0.1)
```